

Computational Intelligence, SS11

<http://www.spsc.tugraz.at/courses/computational-intelligence>

Paul Meissner ¹

6 Homework: Expectation Maximization and K-means Algorithm

[Points: 23+8* , Issued: 2011/06/22 , Deadline: 2011/07/01 , Tutor: Georg Kapeller ²; Info-hour: TBA , Room and Date will be posted in the newsgroup , ;

General: In this homework, you have to implement both the EM and K-means algorithm. For the evaluation, use the dataset provided on the course homepage (`vowels.mat`). The variable `X` contains the unlabeled training data, which is two-dimensional. You already know this dataset from the problem classes.

6.1 Expectation Maximization Algorithm [13 Points]

We want to find a maximum-likelihood Gaussian Mixture Model (GMM) for the training data in X . Using M Gaussian components, this model is given as

$$p(\mathbf{x}|\Theta) = \sum_{m=1}^M \alpha_m \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}_m, \boldsymbol{\Sigma}_m) \quad (1)$$

with the parameter vector Θ containing the component weights α_m , the means $\boldsymbol{\mu}_m$ and the covariance matrices $\boldsymbol{\Sigma}_m$, respectively. The EM algorithm for GMMs tries to iteratively find this ML estimator, i.e., at each iteration i it computes an updated parameter vector Θ_i .

Implement the EM algorithm for GMMs using the following steps:

- 1. Init:** At $i = 0$, select an initial guess of the parameter vector Θ_0 .
- 2. Expectation Step:** For each sample \mathbf{x}_n , calculate the probability $p(m|\mathbf{x}_n, \Theta_i)$ that this sample was caused by the m -th component of the GMM. In the lecture, this value was called r_m^n and is written as

$$r_m^n = \frac{\alpha_{m,i} \mathcal{N}(\mathbf{x}_n|\boldsymbol{\mu}_{m,i}, \boldsymbol{\Sigma}_{m,i})}{\sum_{m'=1}^M \alpha_{m',i} \mathcal{N}(\mathbf{x}_n|\boldsymbol{\mu}_{m',i}, \boldsymbol{\Sigma}_{m',i})} \quad (2)$$

Remember, this corresponds to a Bayesian “soft classification” of each sample.

- 3. Maximization Step:** Using the r_m^n , the effective number of samples for the m -th component is given by $N_m = \sum_{n=1}^N r_m^n$. With this, the entries of Θ can be updated:

$$\boldsymbol{\mu}_{m,i+1} = \frac{1}{N_m} \sum_{n=1}^N r_m^n \mathbf{x}_n \quad (3)$$

$$\boldsymbol{\Sigma}_{m,i+1} = \frac{1}{N_m} \sum_{n=1}^N r_m^n (\mathbf{x}_n - \boldsymbol{\mu}_{m,i+1})(\mathbf{x}_n - \boldsymbol{\mu}_{m,i+1})^T \quad (4)$$

$$\alpha_{m,i+1} = \frac{N_m}{N} \quad (5)$$

¹<mailto:paul.meissner@tugraz.at>

²<mailto:georg.kapeller@student.tugraz.at>

- 4. Likelihood calculation:** Compute the current value of the log-likelihood function of the samples in X , given the current model Θ_{i+1} :

$$\log p(X|\Theta_{i+1}) = \sum_{n=1}^N \log \sum_{m=1}^M \alpha_{m,i+1} \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_{m,i+1}, \boldsymbol{\Sigma}_{m,i+1}) \quad (6)$$

Check whether the likelihood has already converged or not. If yes, terminate the algorithm, if no, go back to step 2.

Your tasks are as follows:

1. Write a Matlab function `[alpha, mu, Sigma, L] = EM(X, M, alpha_0, mu_0, Sigma_0, max_iter)` that implements the EM algorithm! X is the training data, M is the number of Gaussian components, α_0 , μ_0 and Σ_0 are the initial guesses of the parameter vector and `max_iter` is the maximum number of iterations. The function returns the final parameter vector given by `alpha`, `mu` and `Sigma` as well as the log-likelihood over the iterations in `L`. You may use the Matlab function `mvnpdf`.
2. Test your implementation using the dataset given in X . You can compare the result with the labeled data set given in the variables `a`, `e,i,o` and `y`. Make a scatter plot of the data and plot the Gaussian mixture model over this plot. You can use the provided function `plotGaussianContour()` for plotting each of the Gaussian components (the factors α_m are then neglected of course).
3. For your tests, select the correct number of components ($M = 5$), but also check the result when you use more or less components. How do you choose your initialization Θ_0 ? Does this choice have an influence on the result?
4. Also plot the log-likelihood function over the iterations! What is the behavior of this function over the iterations?
5. Within your EM-function, confine the structure of the covariance matrices to diagonal matrices! What is the influence on the result?
6. Make a scatter plot of the data that shows the result of the soft-classification that is done in the E-step. This means, classify each sample using the r_m^n , and plot the samples in different colors for each of the M classes.

6.2 K-means algorithm [10 Points]

In this task, you should implement the K-means algorithm as discussed in the lecture and compare the results with the ones obtained with the EM-algorithm. Remember that you can interpret K-means as a version of the EM-algorithm with several simplifications: First, the classes are just represented by their means, second, there are no class weights and third, a hard classification is performed for all samples. The latter also means that for the parameter update, only the points classified to this component play a role.

1. Write a Matlab function `[mu, D] = k_means(X, M, mu_0, max_iter)` that implements the K-means algorithm. Here, X is the data, M is the number of clusters, μ_0 are the

initial means (cluster centers) and `max_iter` is the maximum number of iterations. The function returns the optimized cluster centers `mu` and the cumulative distance over the iterations in `D`.

2. Perform the same tasks as for the EM-algorithm to evaluate the performance! Of course, the way to plot the classes/components is different now: In the scatter plot, plot the mean value for each class and plot the points that were classified to this class in a certain color.
3. What is the nature of the boundaries between the classes? Compare with the results of the soft-classification in the EM-algorithm! Also compare with the labeled data, can K-means find the class structure well?

6.3 Samples from a Gaussian Mixture Model [8* Points]

1. Write a Matlab function `Y = sampleGMM(alpha, mu, Sigma, N)`, that draws `N` samples from a two-dimensional Gaussian Mixture distribution given by the parameters `alpha`, `mu` and `Sigma`! Hint: Can you split the problem into two parts (First, sample from a specific distribution, then, conditioned on the results from the first step, sample from another one)? You may use the provided function `Y = sampleDiscretePMF(X, PM, N)` that draws `N` samples from a discrete probability distribution `PM`, defined over the values `X`.
2. Using a GMM of your choice ($M > 3$), demonstrate the correctness of your function!